

ANUBHAV SINGH

AI Engineer | Generative AI | LLM & RAG | Agentic AI | vLLM | Voice AI | MLOps
anubhavchauhan.29@gmail.com | +91-8219256891 | [LinkedIn](#) | [Scaler Profile](#) | [Website](#)

AI Engineer specializing in production-grade Generative AI, RAG pipelines, Agentic AI workflows, vLLM model serving (Qwen 3 14B AWQ), Voice AI microservices (Kokoro TTS, Whisper STT), and cloud-native AI on Azure. Experienced with LangChain, LangGraph, Pinecone, FAISS, Docker, and end-to-end MLOps.

TECHNICAL SKILLS

Generative AI & LLMs: LangChain · LangGraph · OpenAI API · RAG · Agentic AI · vLLM · Qwen 3 14B AWQ · Gemma · Prompt Engineering · RAGAS · Self-hosted LLM

Voice AI & Search: Kokoro TTS · OpenAI Whisper STT/ASR · Pinecone · FAISS · Sentence Transformers · Semantic Search · Cross-Encoder Re-ranking

Backend & Cloud: Python · FastAPI · REST API · Microservices · Docker · Microsoft Azure · Azure DevOps · CI/CD · PostgreSQL · MongoDB · MLflow

PROFESSIONAL EXPERIENCE

AI Engineer — Infinity Soft System Aug 2025 – Present

- Designed and deployed production-grade Generative AI & RAG applications using Python, FastAPI, LangChain, LangGraph, and Azure; built end-to-end pipelines covering ingestion, chunking, embedding, vector retrieval, re-ranking, and LLM response synthesis.
- Self-hosted Qwen 3 14B AWQ via vLLM with OpenAI-compatible REST endpoints; deployed Kokoro TTS and Whisper STT microservices to enable a fully local STT → LLM → TTS voice AI pipeline.
- Architected multi-step Agentic AI workflows (LangGraph) for autonomous reasoning, candidate evaluation, and recommendation generation; evaluated pipeline quality with RAGAS metrics (faithfulness, relevancy, context precision/recall).
- Containerized AI microservices with Docker; automated CI/CD via Azure DevOps; integrated Pinecone & FAISS vector databases with Sentence Transformers for high-accuracy semantic search at scale.

KEY PROJECTS

AI-Powered Resume-JD Matching System | Python · FastAPI · LangGraph · Pinecone · FAISS · Gemma · Qwen · Docker · Azure

- Built production RAG platform for semantic resume-JD matching; designed LangGraph Agentic workflows for parsing, skill extraction, scoring, and explainable hiring recommendations; self-hosted Gemma & Qwen for privacy-preserving inference.
- Implemented dense embedding retrieval (Pinecone, FAISS) with cross-encoder re-ranking and deployed scalable FastAPI microservices on Azure with PostgreSQL-backed candidate management.

vLLM LLM & Voice AI Microservices Platform | vLLM · Qwen 3 14B AWQ · Kokoro TTS · Whisper STT · FastAPI · Docker · Azure

- Deployed Qwen 3 14B AWQ via vLLM for high-throughput LLM serving; built Kokoro TTS streaming microservice and Whisper ASR service; wired into end-to-end Voice AI pipeline with GPU-aware Docker deployment on Azure.

RAG Enterprise Knowledge Assistant | Python · FastAPI · LangChain · FAISS · Docker · Azure

- Developed citation-backed conversational AI over enterprise PDFs using FAISS semantic retrieval; reduced hallucinations via retrieval-grounded prompting; deployed production FastAPI services for ingestion and real-time querying.

Water pH Forecasting & MLOps Pipeline | Python · XGBoost · ARIMA · SARIMA · MLflow · FastAPI · Docker · Azure

- Performed comprehensive EDA on water quality data — missing value imputation, outlier detection (IQR/Z-score), distribution analysis, and feature correlation heatmaps — to identify key pH predictors before modelling.
- Benchmarked XGBoost against ARIMA/SARIMA forecasting models; tracked experiments with MLflow; deployed real-time FastAPI inference service with automated retraining pipeline on Azure.

EDUCATION

AI & Machine Learning — Scaler School of Technology 2025

B.Tech, Computer Science & Engineering — Chandigarh University 2022

AREAS OF EXPERTISE

Generative AI · LLM · RAG · Agentic AI · vLLM · AWQ Quantization · Qwen · Gemma · Kokoro TTS · OpenAI Whisper · STT · ASR · Voice AI · Multimodal AI · LangChain · LangGraph · Pinecone · FAISS · Semantic Search · NLP · Prompt Engineering · Sentence Transformers · Python · FastAPI · Microservices · Docker · Azure · CI/CD · MLOps · MLflow · XGBoost · PostgreSQL · Responsible AI · Explainable AI · GPU Inference · Streaming Inference · Cloud-Native AI